

TEJ PRATAP

Data Engineer | Data Analyst | Python | SQL | Databricks | Power BI | M.Sc. AI & Data Science
tej.pratap1227@gmail.com | +49 15510821965 | Munich, Germany | github.com/tejfaster |
linkedin.com/in/tej-pratap-25527a185

PROFESSIONAL SUMMARY

Data professional with 3+ years of experience across analytics engineering, data analysis, and backend development, now pursuing an M.Sc. in Artificial Intelligence and Data Science at TH Deggendorf. **Databricks Certified Data Engineer Professional** and **SQL (Advanced) certified on HackerRank**, with hands-on experience spanning ETL/ELT pipelines, Spark/PySpark, dbt, Delta Lake, Azure Databricks, AWS, PostgreSQL, and Power BI/Tableau reporting. Comfortable owning a problem end-to-end, from raw data to a production pipeline or client-ready dashboard. Currently preparing for the Microsoft Fabric Data Engineer Associate (DP-700) and Databricks Certified Generative AI Engineer Associate certifications.

TECHNICAL SKILLS

Data Engineering: ETL/ELT Pipeline Design, Apache Spark (PySpark), dbt, Data Lake Architecture (Bronze/Silver/Gold), Airflow, Delta Lake, Parquet, Kafka, Data Quality & Validation

Programming & Data: Python (pandas, numpy, scikit-learn), SQL (window functions, CTEs, query optimisation, indexing), R, basic JavaScript, Git, Object-Oriented Programming, API Integration

Cloud & Platforms: Azure Databricks (**Databricks Certified Data Engineer Professional**), AWS (S3, EC2, CloudFront), Docker, Kubernetes (Minikube)

Databases & Modelling: PostgreSQL, Data Warehouse Design, Dimensional Modelling, Star Schema, Snowflake Schema, Enterprise/Relational Data Modelling

Analytics & BI: Power BI (DAX, dashboards, automated reports), Tableau, Excel, PowerPoint, Exploratory Data Analysis (EDA), Statistical Analysis

AI & Machine Learning: Machine Learning, Feature Engineering, CNN/Vision Transformer and Transformer models, Generative AI (certification in progress), LLM Fine-Tuning / Embeddings (in progress)

PROFESSIONAL EXPERIENCE

Analytics Engineer

Dec 2023 – Feb 2025

Autobot, India

- Designed and maintained ETL pipelines integrating APIs, spreadsheets, and internal databases, reducing manual data collection time by 65%.
- Performed data cleaning, transformation, and validation across multiple data sources, increasing reporting accuracy from 82% to 96%.
- Wrote optimised SQL queries for complex reporting needs, improving dashboard query performance by 40%.
- Automated recurring data workflows using Python, saving 20+ hours per month in manual reporting tasks.
- Developed interactive Power BI dashboards and automated reports, improving leadership reporting efficiency by 50% and enabling real-time KPI visibility across the organisation.

Data Analyst & Backend Developer

Sep 2022 – Nov 2023

Antino Labs Pvt. Ltd., India

- Designed relational data models and optimised SQL query paths supporting low-latency analytics at scale on a rapidly growing edtech platform.
- Built AWS cloud infrastructure (S3, EC2, CloudFront) supporting platform data storage and delivery pipelines.
- Integrated behavioural event tracking pipeline (WebEngage) across the app, enabling customer segmentation, lifecycle analytics, and marketing automation for a 500,000+ user platform.
- Designed and maintained Power BI dashboards connecting PostgreSQL data to business reporting views for non-technical stakeholders.

Junior Developer

Jun 2021 – Aug 2022

QuikieBrain LLP, India

- Supported backend data workflows, API documentation, and database schema design in a startup environment.

PROJECTS

CoreSync – Distributed Analytics Pipeline on Real Financial Data

Python, Apache Spark, SQL, Parquet

- Built an end-to-end distributed pipeline processing real Yahoo Finance stock data across USA and EU regions, with multi-region parallel execution and a custom resource-allocation scheduler.
- Designed Bronze/Silver/Gold architecture with data quality validation and offset tracking, ensuring zero duplicate records across pipeline restarts.
- Produced query-optimised Parquet Gold-layer outputs structured for downstream BI and analytics consumption.

MarketPulse – Stock Market Analytics Pipeline

Power BI, SQL, PostgreSQL, Python, dbt, Airflow

- Modelled Gold-layer production tables in PostgreSQL with dbt transformations, end-to-end from raw ingestion to analyst-ready views across 17 stock symbols.

- Delivered a 5-page interactive Power BI dashboard translating technical indicators (RSI, MACD, Bollinger Bands, EMA) into market intelligence, with drill-through filtering by symbol, time range, and signal type.
- Orchestrated pipeline runs with Airflow on a 2-node home cluster.

Retail Analytics – E-Commerce Data Pipeline & BI

Power BI, SQL, PostgreSQL, Python, Star Schema

- Processed 100,000+ real orders from the Olist Brazilian E-Commerce dataset; designed a Star/Snowflake schema (2 fact tables, 6 dimension tables).
- Performed EDA and data validation across all dimensions prior to dashboard build, ensuring analytical accuracy.
- Delivered 5 Power BI dashboards covering sales performance, category trends, payment behaviour, and delivery KPIs; also deployed on Kubernetes/Minikube.

EchoLens – Image Captioning

CNN, Vision Transformer, Transformer (team project)

- Contributed to an image captioning model combining a CNN/Vision Transformer encoder with a Transformer decoder.

EDUCATION

M.Sc. Artificial Intelligence and Data Science

Mar 2025 – Present

TH Deggendorf & University of South Bohemia, Germany

Relevant: Advanced Data Storage & Analysis, Feature Engineering, Advanced Machine Learning, Parallel Computing, Mathematics for AI, Information Theory.

University internship project: Fine-tuning of LLMs for embedding generation (supervisor: Jiří Jelínek).

BCA in Data Science

2018 – 2021

Ajeenkya D Y Patil University, India

Relevant: SQL, Python Programming, Data Mining, Business Analytics, EDA, Big Data, R Programming.

CERTIFICATIONS

Databricks Certified Data Engineer Professional | SQL (Advanced) – HackerRank (previously SQL Basic) | Python for Data Science – IBM | DELL EMC Academic Associate – Data Science | Microsoft Fabric Data Engineer Associate (DP-700) – scheduled Oct 31, 2026 | Databricks Certified Generative AI Engineer Associate – in progress

LANGUAGES

English (fluent) | German (A2) | Hindi (native)